

Keras -- MLPs on MNIST

In [1]:

```
%tensorflow_version 1.9
# if your keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use t
his command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

`%tensorflow_version` only switches the major version: 1.x or 2.x.
You set: `1.9`. This will be interpreted as: `1.x`.

TensorFlow 1.x selected.

Using TensorFlow backend.

In [0]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
%matplotlib inline
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [0]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%
d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%
d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector  
# for each image we have a (28*28) vector  
# we will convert the (28*28) vector into single dimensional vector of 1 * 784  
  
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])  
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors  
  
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%  
d)"%(X_train.shape[1]))  
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d  
)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

In [7]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

In [0]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# example data point after normlizing  
print(X_train[0])
```

5/99

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

```

0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
]

```

In [10]:

```

# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```


In [0]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

In [13]:

```
# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

In [14]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From /tensorflow-1.15.2/python3.6/tensorflow_core/python/ops/math_grad.py:1424: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:207: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:216: The name tf.is_variable_initialized is deprecated. Please use tf.compat.v1.is_variable_initialized instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name tf.variables_initializer is deprecated. Please use tf.compat.v1.variables_initializer instead.

60000/60000 [=====] - 3s 42us/step - loss: 1.2616
- acc: 0.7113 - val_loss: 0.7992 - val_acc: 0.8424

Epoch 2/20

60000/60000 [=====] - 1s 25us/step - loss: 0.7065
- acc: 0.8454 - val_loss: 0.6013 - val_acc: 0.8665

Epoch 3/20

60000/60000 [=====] - 2s 25us/step - loss: 0.5810
- acc: 0.8638 - val_loss: 0.5215 - val_acc: 0.8770

Epoch 4/20

60000/60000 [=====] - 2s 25us/step - loss: 0.5208
- acc: 0.8727 - val_loss: 0.4771 - val_acc: 0.8828

Epoch 5/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4840
- acc: 0.8782 - val_loss: 0.4482 - val_acc: 0.8873

Epoch 6/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4588
- acc: 0.8816 - val_loss: 0.4274 - val_acc: 0.8911

Epoch 7/20
60000/60000 [=====] - 1s 25us/step - loss: 0.4400
- acc: 0.8849 - val_loss: 0.4114 - val_acc: 0.8943

Epoch 8/20
60000/60000 [=====] - 1s 25us/step - loss: 0.4254
- acc: 0.8870 - val_loss: 0.3987 - val_acc: 0.8957

Epoch 9/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4137
- acc: 0.8896 - val_loss: 0.3886 - val_acc: 0.8977

Epoch 10/20
60000/60000 [=====] - 1s 25us/step - loss: 0.4039
- acc: 0.8909 - val_loss: 0.3801 - val_acc: 0.8997

Epoch 11/20
60000/60000 [=====] - 1s 25us/step - loss: 0.3955
- acc: 0.8927 - val_loss: 0.3728 - val_acc: 0.9007

Epoch 12/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3884
- acc: 0.8943 - val_loss: 0.3664 - val_acc: 0.9019

Epoch 13/20
60000/60000 [=====] - 1s 25us/step - loss: 0.3821
- acc: 0.8954 - val_loss: 0.3607 - val_acc: 0.9035

Epoch 14/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3766
- acc: 0.8968 - val_loss: 0.3561 - val_acc: 0.9042

Epoch 15/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3717
- acc: 0.8978 - val_loss: 0.3517 - val_acc: 0.9044

Epoch 16/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3672
- acc: 0.8988 - val_loss: 0.3475 - val_acc: 0.9057

Epoch 17/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3632
- acc: 0.8996 - val_loss: 0.3442 - val_acc: 0.9066

Epoch 18/20
60000/60000 [=====] - 2s 26us/step - loss: 0.3595
- acc: 0.9003 - val_loss: 0.3415 - val_acc: 0.9072

Epoch 19/20
60000/60000 [=====] - 2s 27us/step - loss: 0.3561
- acc: 0.9010 - val_loss: 0.3381 - val_acc: 0.9083

Epoch 20/20
60000/60000 [=====] - 1s 25us/step - loss: 0.3530
- acc: 0.9020 - val_loss: 0.3354 - val_acc: 0.9087

In [15]:

```
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.3353958689153194

Test accuracy: 0.9087



MLP + Sigmoid activation + SGDOptimizer

In [16]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
=====		
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [17]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 31us/step - loss: 2.2699
- acc: 0.2203 - val_loss: 2.2220 - val_acc: 0.4296

Epoch 2/20

60000/60000 [=====] - 2s 28us/step - loss: 2.1776
- acc: 0.4451 - val_loss: 2.1204 - val_acc: 0.5472

Epoch 3/20

60000/60000 [=====] - 2s 27us/step - loss: 2.0578
- acc: 0.5688 - val_loss: 1.9739 - val_acc: 0.6192

Epoch 4/20

60000/60000 [=====] - 2s 26us/step - loss: 1.8843
- acc: 0.6262 - val_loss: 1.7666 - val_acc: 0.6545

Epoch 5/20

60000/60000 [=====] - 2s 27us/step - loss: 1.6600
- acc: 0.6701 - val_loss: 1.5255 - val_acc: 0.7121

Epoch 6/20

60000/60000 [=====] - 2s 27us/step - loss: 1.4250
- acc: 0.7124 - val_loss: 1.2996 - val_acc: 0.7272

Epoch 7/20

60000/60000 [=====] - 2s 26us/step - loss: 1.2214
- acc: 0.7447 - val_loss: 1.1173 - val_acc: 0.7667

Epoch 8/20

60000/60000 [=====] - 2s 27us/step - loss: 1.0620
- acc: 0.7711 - val_loss: 0.9789 - val_acc: 0.7865

Epoch 9/20

60000/60000 [=====] - 2s 27us/step - loss: 0.9404
- acc: 0.7911 - val_loss: 0.8718 - val_acc: 0.8033

Epoch 10/20

60000/60000 [=====] - 2s 26us/step - loss: 0.8467
- acc: 0.8069 - val_loss: 0.7894 - val_acc: 0.8172

Epoch 11/20

60000/60000 [=====] - 2s 27us/step - loss: 0.7731
- acc: 0.8193 - val_loss: 0.7247 - val_acc: 0.8274

Epoch 12/20

60000/60000 [=====] - 2s 27us/step - loss: 0.7141
- acc: 0.8287 - val_loss: 0.6719 - val_acc: 0.8391

Epoch 13/20

60000/60000 [=====] - 2s 27us/step - loss: 0.6663
- acc: 0.8372 - val_loss: 0.6284 - val_acc: 0.8455

Epoch 14/20

60000/60000 [=====] - 2s 26us/step - loss: 0.6269
- acc: 0.8447 - val_loss: 0.5925 - val_acc: 0.8504

Epoch 15/20

60000/60000 [=====] - 2s 27us/step - loss: 0.5940
- acc: 0.8502 - val_loss: 0.5635 - val_acc: 0.8569

Epoch 16/20

60000/60000 [=====] - 2s 26us/step - loss: 0.5664
- acc: 0.8552 - val_loss: 0.5378 - val_acc: 0.8632

Epoch 17/20

60000/60000 [=====] - 2s 27us/step - loss: 0.5429
- acc: 0.8597 - val_loss: 0.5158 - val_acc: 0.8664

Epoch 18/20

60000/60000 [=====] - 2s 27us/step - loss: 0.5225
- acc: 0.8641 - val_loss: 0.4970 - val_acc: 0.8707

Epoch 19/20

60000/60000 [=====] - 2s 27us/step - loss: 0.5048
- acc: 0.8673 - val_loss: 0.4806 - val_acc: 0.8739

Epoch 20/20

60000/60000 [=====] - 2s 27us/step - loss: 0.4890
- acc: 0.8700 - val_loss: 0.4663 - val_acc: 0.8771

In [18]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.46628043448925016

Test accuracy: 0.8771



In [19]:


```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + ADAM

In [20]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290

Total params: 468,874

Trainable params: 468,874

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 36us/step - loss: 0.5341
- acc: 0.8603 - val_loss: 0.2566 - val_acc: 0.9235

Epoch 2/20

60000/60000 [=====] - 2s 31us/step - loss: 0.2208
- acc: 0.9350 - val_loss: 0.1828 - val_acc: 0.9462

Epoch 3/20

60000/60000 [=====] - 2s 31us/step - loss: 0.1632
- acc: 0.9517 - val_loss: 0.1437 - val_acc: 0.9574

Epoch 4/20

60000/60000 [=====] - 2s 30us/step - loss: 0.1266
- acc: 0.9633 - val_loss: 0.1219 - val_acc: 0.9623

Epoch 5/20

60000/60000 [=====] - 2s 31us/step - loss: 0.1001
- acc: 0.9706 - val_loss: 0.0997 - val_acc: 0.9697

Epoch 6/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0799
- acc: 0.9767 - val_loss: 0.0938 - val_acc: 0.9708

Epoch 7/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0651
- acc: 0.9799 - val_loss: 0.0816 - val_acc: 0.9747

Epoch 8/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0529
- acc: 0.9845 - val_loss: 0.0772 - val_acc: 0.9770

Epoch 9/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0422
- acc: 0.9881 - val_loss: 0.0754 - val_acc: 0.9764

Epoch 10/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0344
- acc: 0.9900 - val_loss: 0.0663 - val_acc: 0.9801

Epoch 11/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0278
- acc: 0.9924 - val_loss: 0.0631 - val_acc: 0.9794

Epoch 12/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0225
- acc: 0.9940 - val_loss: 0.0624 - val_acc: 0.9816

Epoch 13/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0180
- acc: 0.9955 - val_loss: 0.0632 - val_acc: 0.9802

Epoch 14/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0148
- acc: 0.9963 - val_loss: 0.0646 - val_acc: 0.9811

Epoch 15/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0115
- acc: 0.9973 - val_loss: 0.0606 - val_acc: 0.9824

Epoch 16/20

```
60000/60000 [=====] - 2s 31us/step - loss: 0.0102
- acc: 0.9977 - val_loss: 0.0640 - val_acc: 0.9818
Epoch 17/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0075
- acc: 0.9985 - val_loss: 0.0674 - val_acc: 0.9808
Epoch 18/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0062
- acc: 0.9987 - val_loss: 0.0631 - val_acc: 0.9828
Epoch 19/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0042
- acc: 0.9993 - val_loss: 0.0706 - val_acc: 0.9806
Epoch 20/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0040
- acc: 0.9990 - val_loss: 0.0753 - val_acc: 0.9806
```

In [21]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07527304192187294
Test accuracy: 0.9806



In [22]:

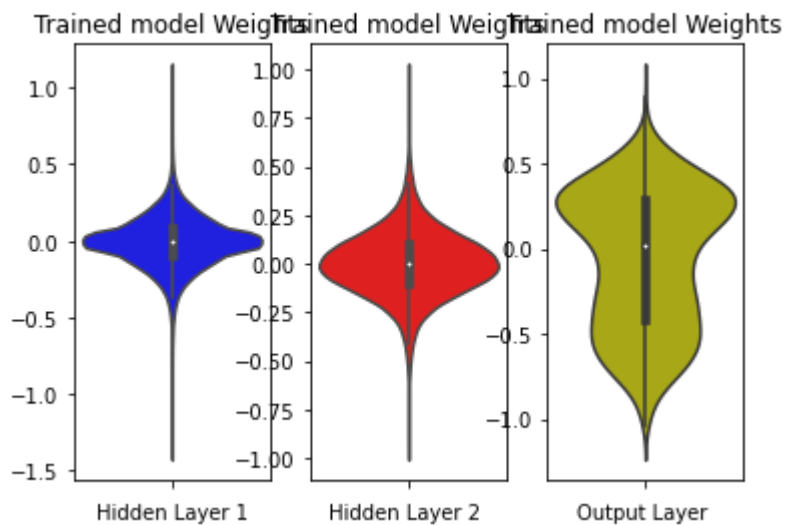
```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU +SGD

In [23]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(\theta, \sigma) = N(\theta, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(\theta, \sigma) = N(\theta, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4409: The name tf.random_normal is deprecated. Please use tf.random.normal instead.

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [24]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 32us/step - loss: 0.7488
- acc: 0.7907 - val_loss: 0.3851 - val_acc: 0.8964

Epoch 2/20

60000/60000 [=====] - 2s 26us/step - loss: 0.3489
- acc: 0.9021 - val_loss: 0.2973 - val_acc: 0.9160

Epoch 3/20

60000/60000 [=====] - 2s 26us/step - loss: 0.2862
- acc: 0.9180 - val_loss: 0.2582 - val_acc: 0.9274

Epoch 4/20

60000/60000 [=====] - 2s 27us/step - loss: 0.2520
- acc: 0.9283 - val_loss: 0.2341 - val_acc: 0.9334

Epoch 5/20

60000/60000 [=====] - 2s 26us/step - loss: 0.2282
- acc: 0.9350 - val_loss: 0.2156 - val_acc: 0.9371

Epoch 6/20

60000/60000 [=====] - 2s 27us/step - loss: 0.2105
- acc: 0.9404 - val_loss: 0.2025 - val_acc: 0.9417

Epoch 7/20

60000/60000 [=====] - 2s 26us/step - loss: 0.1958
- acc: 0.9440 - val_loss: 0.1898 - val_acc: 0.9457

Epoch 8/20

60000/60000 [=====] - 2s 25us/step - loss: 0.1835
- acc: 0.9479 - val_loss: 0.1814 - val_acc: 0.9460

Epoch 9/20

60000/60000 [=====] - 2s 26us/step - loss: 0.1730
- acc: 0.9511 - val_loss: 0.1713 - val_acc: 0.9492

Epoch 10/20

60000/60000 [=====] - 2s 27us/step - loss: 0.1638
- acc: 0.9532 - val_loss: 0.1640 - val_acc: 0.9512

Epoch 11/20

60000/60000 [=====] - 2s 27us/step - loss: 0.1555
- acc: 0.9556 - val_loss: 0.1598 - val_acc: 0.9530

Epoch 12/20

60000/60000 [=====] - 2s 28us/step - loss: 0.1480
- acc: 0.9580 - val_loss: 0.1518 - val_acc: 0.9552

Epoch 13/20

60000/60000 [=====] - 2s 27us/step - loss: 0.1414
- acc: 0.9597 - val_loss: 0.1462 - val_acc: 0.9568

Epoch 14/20

60000/60000 [=====] - 2s 26us/step - loss: 0.1352
- acc: 0.9618 - val_loss: 0.1442 - val_acc: 0.9580

Epoch 15/20

60000/60000 [=====] - 2s 27us/step - loss: 0.1296
- acc: 0.9634 - val_loss: 0.1379 - val_acc: 0.9592

Epoch 16/20

60000/60000 [=====] - 2s 26us/step - loss: 0.1245
- acc: 0.9651 - val_loss: 0.1351 - val_acc: 0.9604

Epoch 17/20

60000/60000 [=====] - 2s 26us/step - loss: 0.1198
- acc: 0.9665 - val_loss: 0.1304 - val_acc: 0.9604

Epoch 18/20

60000/60000 [=====] - 2s 26us/step - loss: 0.1154
- acc: 0.9679 - val_loss: 0.1261 - val_acc: 0.9617

Epoch 19/20

60000/60000 [=====] - 2s 27us/step - loss: 0.1111
- acc: 0.9689 - val_loss: 0.1229 - val_acc: 0.9624

Epoch 20/20

60000/60000 [=====] - 2s 27us/step - loss: 0.1073
- acc: 0.9702 - val_loss: 0.1221 - val_acc: 0.9627

In [25]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12213720679543913

Test accuracy: 0.9627



In [26]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + ADAM

In [27]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290

Total params: 468,874

Trainable params: 468,874

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 36us/step - loss: 0.2344

- acc: 0.9307 - val_loss: 0.1179 - val_acc: 0.9626

Epoch 2/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0863

- acc: 0.9736 - val_loss: 0.0837 - val_acc: 0.9713

Epoch 3/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0533

- acc: 0.9838 - val_loss: 0.0725 - val_acc: 0.9762

Epoch 4/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0347

- acc: 0.9897 - val_loss: 0.0827 - val_acc: 0.9739

Epoch 5/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0261

- acc: 0.9917 - val_loss: 0.0751 - val_acc: 0.9776

Epoch 6/20

60000/60000 [=====] - 2s 32us/step - loss: 0.0198

- acc: 0.9937 - val_loss: 0.0749 - val_acc: 0.9785

Epoch 7/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0176

- acc: 0.9943 - val_loss: 0.0723 - val_acc: 0.9801

Epoch 8/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0136

- acc: 0.9957 - val_loss: 0.0813 - val_acc: 0.9787

Epoch 9/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0137

- acc: 0.9951 - val_loss: 0.0784 - val_acc: 0.9778

Epoch 10/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0132

- acc: 0.9956 - val_loss: 0.0808 - val_acc: 0.9801

Epoch 11/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0117

- acc: 0.9960 - val_loss: 0.0885 - val_acc: 0.9782

Epoch 12/20

60000/60000 [=====] - 2s 29us/step - loss: 0.0089

- acc: 0.9970 - val_loss: 0.1125 - val_acc: 0.9748

Epoch 13/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0105

- acc: 0.9964 - val_loss: 0.0865 - val_acc: 0.9805

Epoch 14/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0079

- acc: 0.9972 - val_loss: 0.0933 - val_acc: 0.9800

Epoch 15/20

60000/60000 [=====] - 2s 30us/step - loss: 0.0109

- acc: 0.9965 - val_loss: 0.0812 - val_acc: 0.9833

Epoch 16/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0070

- acc: 0.9976 - val_loss: 0.0904 - val_acc: 0.9812

Epoch 17/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0109

- acc: 0.9966 - val_loss: 0.0968 - val_acc: 0.9799

Epoch 18/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0089

- acc: 0.9970 - val_loss: 0.0945 - val_acc: 0.9803

Epoch 19/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0021

- acc: 0.9994 - val_loss: 0.0951 - val_acc: 0.9825

Epoch 20/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0063

- acc: 0.9981 - val_loss: 0.1084 - val_acc: 0.9785

In [28]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10837043899441291
Test accuracy: 0.9785



In [29]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [30]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [31]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 61us/step - loss: 0.3026
- acc: 0.9102 - val_loss: 0.2071 - val_acc: 0.9416

Epoch 2/20

60000/60000 [=====] - 3s 48us/step - loss: 0.1736
- acc: 0.9497 - val_loss: 0.1710 - val_acc: 0.9504

Epoch 3/20

60000/60000 [=====] - 3s 48us/step - loss: 0.1374
- acc: 0.9598 - val_loss: 0.1511 - val_acc: 0.9542

Epoch 4/20

60000/60000 [=====] - 3s 49us/step - loss: 0.1140
- acc: 0.9656 - val_loss: 0.1368 - val_acc: 0.9589

Epoch 5/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0949
- acc: 0.9712 - val_loss: 0.1299 - val_acc: 0.9603

Epoch 6/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0811
- acc: 0.9755 - val_loss: 0.1175 - val_acc: 0.9647

Epoch 7/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0696
- acc: 0.9780 - val_loss: 0.1096 - val_acc: 0.9661

Epoch 8/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0589
- acc: 0.9820 - val_loss: 0.1067 - val_acc: 0.9688

Epoch 9/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0493
- acc: 0.9847 - val_loss: 0.1097 - val_acc: 0.9682

Epoch 10/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0434
- acc: 0.9859 - val_loss: 0.1041 - val_acc: 0.9698

Epoch 11/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0376
- acc: 0.9879 - val_loss: 0.0994 - val_acc: 0.9697

Epoch 12/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0349
- acc: 0.9887 - val_loss: 0.0974 - val_acc: 0.9722

Epoch 13/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0307
- acc: 0.9903 - val_loss: 0.1008 - val_acc: 0.9723

Epoch 14/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0306
- acc: 0.9900 - val_loss: 0.0929 - val_acc: 0.9745

Epoch 15/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0237
- acc: 0.9924 - val_loss: 0.0998 - val_acc: 0.9732

Epoch 16/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0201
- acc: 0.9938 - val_loss: 0.1011 - val_acc: 0.9731

Epoch 17/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0211
- acc: 0.9930 - val_loss: 0.1067 - val_acc: 0.9707

Epoch 18/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0209
- acc: 0.9931 - val_loss: 0.1049 - val_acc: 0.9719

Epoch 19/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0160
- acc: 0.9951 - val_loss: 0.0969 - val_acc: 0.9737

Epoch 20/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0152
- acc: 0.9949 - val_loss: 0.0991 - val_acc: 0.9741

In [32]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09913447910312971

Test accuracy: 0.9741



In [33]:


```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm using 3 hidden Layers + AdamOptimizer

In [34]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(250, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(200, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(170, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 250)	196250
batch_normalization_3 (Batch Normalization)	(None, 250)	1000
dense_18 (Dense)	(None, 200)	50200
batch_normalization_4 (Batch Normalization)	(None, 200)	800
dense_19 (Dense)	(None, 170)	34170
batch_normalization_5 (Batch Normalization)	(None, 170)	680
dense_20 (Dense)	(None, 10)	1710
Total params: 284,810		
Trainable params: 283,570		
Non-trainable params: 1,240		

In [35]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 76us/step - loss: 0.2743

- acc: 0.9182 - val_loss: 0.1835 - val_acc: 0.9450

Epoch 2/20

60000/60000 [=====] - 4s 60us/step - loss: 0.1370

- acc: 0.9581 - val_loss: 0.1298 - val_acc: 0.9609

Epoch 3/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0957

- acc: 0.9707 - val_loss: 0.1306 - val_acc: 0.9613

Epoch 4/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0737

- acc: 0.9761 - val_loss: 0.1282 - val_acc: 0.9619

Epoch 5/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0569

- acc: 0.9813 - val_loss: 0.1324 - val_acc: 0.9610

Epoch 6/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0504

- acc: 0.9826 - val_loss: 0.1231 - val_acc: 0.9630

Epoch 7/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0410

- acc: 0.9864 - val_loss: 0.1235 - val_acc: 0.9670

Epoch 8/20

60000/60000 [=====] - 4s 64us/step - loss: 0.0350

- acc: 0.9883 - val_loss: 0.1077 - val_acc: 0.9689

Epoch 9/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0305

- acc: 0.9900 - val_loss: 0.1013 - val_acc: 0.9726

Epoch 10/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0282

- acc: 0.9907 - val_loss: 0.1127 - val_acc: 0.9695

Epoch 11/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0228

- acc: 0.9924 - val_loss: 0.1147 - val_acc: 0.9703

Epoch 12/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0221

- acc: 0.9926 - val_loss: 0.1099 - val_acc: 0.9719

Epoch 13/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0234

- acc: 0.9920 - val_loss: 0.1087 - val_acc: 0.9704

Epoch 14/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0168

- acc: 0.9942 - val_loss: 0.1093 - val_acc: 0.9729

Epoch 15/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0189

- acc: 0.9936 - val_loss: 0.1127 - val_acc: 0.9726

Epoch 16/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0183

- acc: 0.9938 - val_loss: 0.1163 - val_acc: 0.9706

Epoch 17/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0172

- acc: 0.9941 - val_loss: 0.1140 - val_acc: 0.9736

Epoch 18/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0163

- acc: 0.9942 - val_loss: 0.1151 - val_acc: 0.9703

Epoch 19/20

60000/60000 [=====] - 4s 64us/step - loss: 0.0126

- acc: 0.9958 - val_loss: 0.1122 - val_acc: 0.9757

Epoch 20/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0146

- acc: 0.9948 - val_loss: 0.1231 - val_acc: 0.9736

In [36]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12309521172174281

Test accuracy: 0.9736



In [37]:

```
w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='c')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm using 5 hidden Layers + AdamOptimizer

In [38]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(600, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(500, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(400, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(300, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(200, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_21 (Dense)	(None, 600)	471000
batch_normalization_6 (Batch Normalization)	(None, 600)	2400
dense_22 (Dense)	(None, 500)	300500
batch_normalization_7 (Batch Normalization)	(None, 500)	2000
dense_23 (Dense)	(None, 400)	200400
batch_normalization_8 (Batch Normalization)	(None, 400)	1600
dense_24 (Dense)	(None, 300)	120300
batch_normalization_9 (Batch Normalization)	(None, 300)	1200
dense_25 (Dense)	(None, 200)	60200
batch_normalization_10 (Batch Normalization)	(None, 200)	800
dense_26 (Dense)	(None, 10)	2010
Total params: 1,162,410		
Trainable params: 1,158,410		
Non-trainable params: 4,000		

In [39]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```


Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 114us/step - loss: 0.268

2 - acc: 0.9204 - val_loss: 0.1818 - val_acc: 0.9423

Epoch 2/20

60000/60000 [=====] - 6s 92us/step - loss: 0.1321

- acc: 0.9597 - val_loss: 0.1601 - val_acc: 0.9535

Epoch 3/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1019

- acc: 0.9677 - val_loss: 0.1247 - val_acc: 0.9618

Epoch 4/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0778

- acc: 0.9753 - val_loss: 0.1060 - val_acc: 0.9686

Epoch 5/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0701

- acc: 0.9768 - val_loss: 0.1200 - val_acc: 0.9639

Epoch 6/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0562

- acc: 0.9818 - val_loss: 0.0947 - val_acc: 0.9727

Epoch 7/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0489

- acc: 0.9834 - val_loss: 0.1080 - val_acc: 0.9695

Epoch 8/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0451

- acc: 0.9853 - val_loss: 0.0951 - val_acc: 0.9730

Epoch 9/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0376

- acc: 0.9878 - val_loss: 0.0860 - val_acc: 0.9753

Epoch 10/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0343

- acc: 0.9884 - val_loss: 0.0881 - val_acc: 0.9747

Epoch 11/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0328

- acc: 0.9888 - val_loss: 0.0863 - val_acc: 0.9752

Epoch 12/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0268

- acc: 0.9911 - val_loss: 0.0880 - val_acc: 0.9760

Epoch 13/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0292

- acc: 0.9905 - val_loss: 0.0954 - val_acc: 0.9734

Epoch 14/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0245

- acc: 0.9920 - val_loss: 0.0824 - val_acc: 0.9768

Epoch 15/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0246

- acc: 0.9920 - val_loss: 0.0963 - val_acc: 0.9762

Epoch 16/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0245

- acc: 0.9917 - val_loss: 0.0847 - val_acc: 0.9776

Epoch 17/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0206

- acc: 0.9927 - val_loss: 0.0960 - val_acc: 0.9770

Epoch 18/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0242

- acc: 0.9920 - val_loss: 0.0991 - val_acc: 0.9758

Epoch 19/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0179

- acc: 0.9937 - val_loss: 0.0971 - val_acc: 0.9759

Epoch 20/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0183

- acc: 0.9935 - val_loss: 0.0959 - val_acc: 0.9766

In [40]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09593826632477576

Test accuracy: 0.9766



In [41]:

```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[2].flatten().reshape(-1,1)
h4_w = w_after[2].flatten().reshape(-1,1)
h5_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='c')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='g')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='m')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm using 7 hidden Layers + AdamOptimizer

In [42]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(720, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(650, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(580, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(510, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(450, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(390, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_27 (Dense)	(None, 512)	401920
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
dense_28 (Dense)	(None, 720)	369360
batch_normalization_12 (Batch Normalization)	(None, 720)	2880
dense_29 (Dense)	(None, 650)	468650
batch_normalization_13 (Batch Normalization)	(None, 650)	2600
dense_30 (Dense)	(None, 580)	377580
batch_normalization_14 (Batch Normalization)	(None, 580)	2320
dense_31 (Dense)	(None, 510)	296310
batch_normalization_15 (Batch Normalization)	(None, 510)	2040
dense_32 (Dense)	(None, 450)	229950
batch_normalization_16 (Batch Normalization)	(None, 450)	1800
dense_33 (Dense)	(None, 390)	175890
batch_normalization_17 (Batch Normalization)	(None, 390)	1560
dense_34 (Dense)	(None, 10)	3910
Total params: 2,338,818		
Trainable params: 2,331,194		
Non-trainable params: 7,624		

In [43]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 150us/step - loss: 0.286

5 - acc: 0.9165 - val_loss: 0.1705 - val_acc: 0.9518

Epoch 2/20

60000/60000 [=====] - 7s 113us/step - loss: 0.151

3 - acc: 0.9553 - val_loss: 0.1444 - val_acc: 0.9573

Epoch 3/20

60000/60000 [=====] - 7s 112us/step - loss: 0.111

5 - acc: 0.9661 - val_loss: 0.1376 - val_acc: 0.9631

Epoch 4/20

60000/60000 [=====] - 7s 115us/step - loss: 0.094

7 - acc: 0.9707 - val_loss: 0.1137 - val_acc: 0.9671

Epoch 5/20

60000/60000 [=====] - 7s 113us/step - loss: 0.071

9 - acc: 0.9779 - val_loss: 0.1100 - val_acc: 0.9680

Epoch 6/20

60000/60000 [=====] - 7s 111us/step - loss: 0.065

6 - acc: 0.9794 - val_loss: 0.1012 - val_acc: 0.9720

Epoch 7/20

60000/60000 [=====] - 7s 117us/step - loss: 0.056

6 - acc: 0.9824 - val_loss: 0.1200 - val_acc: 0.9682

Epoch 8/20

60000/60000 [=====] - 7s 113us/step - loss: 0.050

4 - acc: 0.9843 - val_loss: 0.1058 - val_acc: 0.9723

Epoch 9/20

60000/60000 [=====] - 7s 109us/step - loss: 0.046

0 - acc: 0.9856 - val_loss: 0.1037 - val_acc: 0.9725

Epoch 10/20

60000/60000 [=====] - 7s 109us/step - loss: 0.041

6 - acc: 0.9868 - val_loss: 0.1276 - val_acc: 0.9667

Epoch 11/20

60000/60000 [=====] - 7s 109us/step - loss: 0.038

3 - acc: 0.9882 - val_loss: 0.1126 - val_acc: 0.9716

Epoch 12/20

60000/60000 [=====] - 7s 109us/step - loss: 0.039

6 - acc: 0.9883 - val_loss: 0.1085 - val_acc: 0.9718

Epoch 13/20

60000/60000 [=====] - 7s 112us/step - loss: 0.032

1 - acc: 0.9895 - val_loss: 0.0917 - val_acc: 0.9754

Epoch 14/20

60000/60000 [=====] - 7s 112us/step - loss: 0.030

6 - acc: 0.9905 - val_loss: 0.1089 - val_acc: 0.9722

Epoch 15/20

60000/60000 [=====] - 7s 111us/step - loss: 0.032

1 - acc: 0.9896 - val_loss: 0.1152 - val_acc: 0.9716

Epoch 16/20

60000/60000 [=====] - 7s 114us/step - loss: 0.027

6 - acc: 0.9912 - val_loss: 0.1003 - val_acc: 0.9751

Epoch 17/20

60000/60000 [=====] - 7s 110us/step - loss: 0.024

4 - acc: 0.9926 - val_loss: 0.1074 - val_acc: 0.9739

Epoch 18/20

60000/60000 [=====] - 7s 111us/step - loss: 0.025

6 - acc: 0.9915 - val_loss: 0.1112 - val_acc: 0.9739

Epoch 19/20

60000/60000 [=====] - 7s 112us/step - loss: 0.020

4 - acc: 0.9937 - val_loss: 0.1008 - val_acc: 0.9770

Epoch 20/20

60000/60000 [=====] - 7s 111us/step - loss: 0.025

3 - acc: 0.9925 - val_loss: 0.0990 - val_acc: 0.9760

In [44]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09895186286373064

Test accuracy: 0.976



In [45]:

```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[4].flatten().reshape(-1,1)
h5_w = w_after[4].flatten().reshape(-1,1)
h6_w = w_after[4].flatten().reshape(-1,1)
h7_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[8].flatten().reshape(-1,1)

fig = plt.figure(figsize =(20,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 8, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 8, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 8, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')


plt.subplot(1, 8, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='c')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 8, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='w')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 8, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h6_w, color='b')
plt.xlabel('Hidden Layer 6 ')

plt.subplot(1, 8, 7)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h7_w, color='m')
plt.xlabel('Hidden Layer 7 ')

plt.subplot(1, 8, 8)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5. MLP + 2 Dropout's + AdamOptimizer

In [46]:

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Model: "sequential_10"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 128)	65664
dropout_2 (Dropout)	(None, 128)	0
dense_37 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [47]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 59us/step - loss: 0.9635
- acc: 0.6934 - val_loss: 0.3132 - val_acc: 0.9074

Epoch 2/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4220
- acc: 0.8755 - val_loss: 0.2321 - val_acc: 0.9292

Epoch 3/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3325
- acc: 0.9022 - val_loss: 0.1930 - val_acc: 0.9415

Epoch 4/20

60000/60000 [=====] - 2s 36us/step - loss: 0.2764
- acc: 0.9184 - val_loss: 0.1652 - val_acc: 0.9481

Epoch 5/20

60000/60000 [=====] - 2s 37us/step - loss: 0.2397
- acc: 0.9291 - val_loss: 0.1425 - val_acc: 0.9555

Epoch 6/20

60000/60000 [=====] - 2s 36us/step - loss: 0.2129
- acc: 0.9382 - val_loss: 0.1291 - val_acc: 0.9598

Epoch 7/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1912
- acc: 0.9437 - val_loss: 0.1180 - val_acc: 0.9628

Epoch 8/20

60000/60000 [=====] - 2s 37us/step - loss: 0.1740
- acc: 0.9498 - val_loss: 0.1067 - val_acc: 0.9651

Epoch 9/20

60000/60000 [=====] - 2s 34us/step - loss: 0.1566
- acc: 0.9541 - val_loss: 0.1030 - val_acc: 0.9678

Epoch 10/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1461
- acc: 0.9571 - val_loss: 0.0938 - val_acc: 0.9708

Epoch 11/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1367
- acc: 0.9596 - val_loss: 0.0880 - val_acc: 0.9722

Epoch 12/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1268
- acc: 0.9619 - val_loss: 0.0862 - val_acc: 0.9727

Epoch 13/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1193
- acc: 0.9642 - val_loss: 0.0863 - val_acc: 0.9745

Epoch 14/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1141
- acc: 0.9657 - val_loss: 0.0787 - val_acc: 0.9762

Epoch 15/20

60000/60000 [=====] - 2s 36us/step - loss: 0.1072
- acc: 0.9681 - val_loss: 0.0785 - val_acc: 0.9764

Epoch 16/20

60000/60000 [=====] - 2s 35us/step - loss: 0.1029
- acc: 0.9693 - val_loss: 0.0754 - val_acc: 0.9775

Epoch 17/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0979
- acc: 0.9711 - val_loss: 0.0743 - val_acc: 0.9772

Epoch 18/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0927
- acc: 0.9720 - val_loss: 0.0739 - val_acc: 0.9773

Epoch 19/20

60000/60000 [=====] - 2s 36us/step - loss: 0.0868
- acc: 0.9738 - val_loss: 0.0749 - val_acc: 0.9777

Epoch 20/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0850
- acc: 0.9741 - val_loss: 0.0698 - val_acc: 0.9789

In [48]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06975141780656995

Test accuracy: 0.9789



In [49]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



1. MLP + 3 Dropout's + AdamOptimizer

In [55]:

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(496, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(Dropout(0.7))

model_drop.add(Dense(279, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

WARNING:tensorflow:Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Model: "sequential_12"

Layer (type)	Output Shape	Param #
dense_42 (Dense)	(None, 512)	401920
dropout_6 (Dropout)	(None, 512)	0
dense_43 (Dense)	(None, 496)	254448
dropout_7 (Dropout)	(None, 496)	0
dense_44 (Dense)	(None, 279)	138663
dropout_8 (Dropout)	(None, 279)	0
dense_45 (Dense)	(None, 10)	2800
Total params: 797,831		
Trainable params: 797,831		
Non-trainable params: 0		

In [56]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```


Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 66us/step - loss: 1.0892
- acc: 0.6844 - val_loss: 0.2835 - val_acc: 0.9236

Epoch 2/20

60000/60000 [=====] - 2s 39us/step - loss: 0.4311
- acc: 0.8779 - val_loss: 0.2002 - val_acc: 0.9486

Epoch 3/20

60000/60000 [=====] - 2s 39us/step - loss: 0.3407
- acc: 0.9052 - val_loss: 0.1691 - val_acc: 0.9541

Epoch 4/20

60000/60000 [=====] - 2s 38us/step - loss: 0.2960
- acc: 0.9193 - val_loss: 0.1497 - val_acc: 0.9589

Epoch 5/20

60000/60000 [=====] - 2s 40us/step - loss: 0.2601
- acc: 0.9295 - val_loss: 0.1421 - val_acc: 0.9587

Epoch 6/20

60000/60000 [=====] - 2s 38us/step - loss: 0.2414
- acc: 0.9342 - val_loss: 0.1290 - val_acc: 0.9637

Epoch 7/20

60000/60000 [=====] - 2s 40us/step - loss: 0.2259
- acc: 0.9382 - val_loss: 0.1269 - val_acc: 0.9641

Epoch 8/20

60000/60000 [=====] - 2s 40us/step - loss: 0.2111
- acc: 0.9426 - val_loss: 0.1115 - val_acc: 0.9670

Epoch 9/20

60000/60000 [=====] - 2s 39us/step - loss: 0.2012
- acc: 0.9452 - val_loss: 0.1162 - val_acc: 0.9647

Epoch 10/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1971
- acc: 0.9460 - val_loss: 0.1165 - val_acc: 0.9668

Epoch 11/20

60000/60000 [=====] - 2s 38us/step - loss: 0.1928
- acc: 0.9484 - val_loss: 0.1159 - val_acc: 0.9673

Epoch 12/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1825
- acc: 0.9514 - val_loss: 0.1104 - val_acc: 0.9709

Epoch 13/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1785
- acc: 0.9520 - val_loss: 0.1085 - val_acc: 0.9699

Epoch 14/20

60000/60000 [=====] - 2s 38us/step - loss: 0.1734
- acc: 0.9539 - val_loss: 0.1016 - val_acc: 0.9726

Epoch 15/20

60000/60000 [=====] - 2s 41us/step - loss: 0.1701
- acc: 0.9539 - val_loss: 0.1045 - val_acc: 0.9715

Epoch 16/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1611
- acc: 0.9581 - val_loss: 0.1061 - val_acc: 0.9714

Epoch 17/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1587
- acc: 0.9571 - val_loss: 0.1027 - val_acc: 0.9712

Epoch 18/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1581
- acc: 0.9574 - val_loss: 0.1015 - val_acc: 0.9741

Epoch 19/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1522
- acc: 0.9591 - val_loss: 0.1037 - val_acc: 0.9738

Epoch 20/20

60000/60000 [=====] - 2s 38us/step - loss: 0.1488
- acc: 0.9605 - val_loss: 0.1049 - val_acc: 0.9747

In [57]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10485141305003781

Test accuracy: 0.9747



In [58]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



1. MLP + Batch-norm + 5 Dropout's + AdamOptimizer

In [60]:

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(420, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(Dropout(0.4))

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.8))

model_drop.add(Dense(638, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.25))

model_drop.add(Dense(381, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.17))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

WARNING:tensorflow:Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Model: "sequential_14"

Layer (type)	Output Shape	Param #
=====		
dense_52 (Dense)	(None, 420)	329700
dropout_14 (Dropout)	(None, 420)	0
dense_53 (Dense)	(None, 128)	53888
dropout_15 (Dropout)	(None, 128)	0
dense_54 (Dense)	(None, 512)	66048
dropout_16 (Dropout)	(None, 512)	0
dense_55 (Dense)	(None, 638)	327294
dropout_17 (Dropout)	(None, 638)	0
dense_56 (Dense)	(None, 381)	243459
dropout_18 (Dropout)	(None, 381)	0
dense_57 (Dense)	(None, 10)	3820
=====		
Total params: 1,024,209		
Trainable params: 1,024,209		
Non-trainable params: 0		

In [61]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 79us/step - loss: 1.0027
- acc: 0.6563 - val_loss: 0.2651 - val_acc: 0.9257

Epoch 2/20

60000/60000 [=====] - 3s 46us/step - loss: 0.4191
- acc: 0.8818 - val_loss: 0.1868 - val_acc: 0.9475

Epoch 3/20

60000/60000 [=====] - 3s 46us/step - loss: 0.3330
- acc: 0.9082 - val_loss: 0.1680 - val_acc: 0.9538

Epoch 4/20

60000/60000 [=====] - 3s 46us/step - loss: 0.2898
- acc: 0.9219 - val_loss: 0.1554 - val_acc: 0.9593

Epoch 5/20

60000/60000 [=====] - 3s 45us/step - loss: 0.2577
- acc: 0.9315 - val_loss: 0.1381 - val_acc: 0.9618

Epoch 6/20

60000/60000 [=====] - 3s 44us/step - loss: 0.2448
- acc: 0.9360 - val_loss: 0.1366 - val_acc: 0.9634

Epoch 7/20

60000/60000 [=====] - 3s 46us/step - loss: 0.2306
- acc: 0.9408 - val_loss: 0.1331 - val_acc: 0.9651

Epoch 8/20

60000/60000 [=====] - 3s 45us/step - loss: 0.2165
- acc: 0.9441 - val_loss: 0.1384 - val_acc: 0.9652

Epoch 9/20

60000/60000 [=====] - 3s 46us/step - loss: 0.2065
- acc: 0.9477 - val_loss: 0.1337 - val_acc: 0.9688

Epoch 10/20

60000/60000 [=====] - 3s 46us/step - loss: 0.1955
- acc: 0.9492 - val_loss: 0.1197 - val_acc: 0.9688

Epoch 11/20

60000/60000 [=====] - 3s 45us/step - loss: 0.1898
- acc: 0.9532 - val_loss: 0.1180 - val_acc: 0.9704

Epoch 12/20

60000/60000 [=====] - 3s 45us/step - loss: 0.1870
- acc: 0.9529 - val_loss: 0.1187 - val_acc: 0.9703

Epoch 13/20

60000/60000 [=====] - 3s 45us/step - loss: 0.1836
- acc: 0.9540 - val_loss: 0.1218 - val_acc: 0.9697

Epoch 14/20

60000/60000 [=====] - 3s 46us/step - loss: 0.1781
- acc: 0.9559 - val_loss: 0.1167 - val_acc: 0.9719

Epoch 15/20

60000/60000 [=====] - 3s 46us/step - loss: 0.1739
- acc: 0.9565 - val_loss: 0.1144 - val_acc: 0.9716

Epoch 16/20

60000/60000 [=====] - 3s 45us/step - loss: 0.1622
- acc: 0.9591 - val_loss: 0.1230 - val_acc: 0.9737

Epoch 17/20

60000/60000 [=====] - 3s 46us/step - loss: 0.1644
- acc: 0.9587 - val_loss: 0.1172 - val_acc: 0.9719

Epoch 18/20

60000/60000 [=====] - 3s 44us/step - loss: 0.1570
- acc: 0.9606 - val_loss: 0.1219 - val_acc: 0.9700

Epoch 19/20

60000/60000 [=====] - 3s 46us/step - loss: 0.1588
- acc: 0.9593 - val_loss: 0.1074 - val_acc: 0.9739

Epoch 20/20

60000/60000 [=====] - 3s 44us/step - loss: 0.1519
- acc: 0.9630 - val_loss: 0.1073 - val_acc: 0.9740

In [62]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10725452537201345

Test accuracy: 0.974



In [63]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[4].flatten().reshape(-1,1)
h5_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure(figsize=(20,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 5, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 5, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 5, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1,5, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='w')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 5, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



6. MLP + Batch-norm + Dropout + AdamOptimizer

In [64]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
=====		
dense_58 (Dense)	(None, 512)	401920
batch_normalization_18 (Batch Normalization)	(None, 512)	2048
dropout_19 (Dropout)	(None, 512)	0
dense_59 (Dense)	(None, 128)	65664
batch_normalization_19 (Batch Normalization)	(None, 128)	512
dropout_20 (Dropout)	(None, 128)	0
dense_60 (Dense)	(None, 10)	1290
=====		
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [65]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 93us/step - loss: 0.6686
- acc: 0.7923 - val_loss: 0.2930 - val_acc: 0.9139

Epoch 2/20

60000/60000 [=====] - 3s 55us/step - loss: 0.4310
- acc: 0.8690 - val_loss: 0.2566 - val_acc: 0.9244

Epoch 3/20

60000/60000 [=====] - 3s 55us/step - loss: 0.3884
- acc: 0.8829 - val_loss: 0.2397 - val_acc: 0.9287

Epoch 4/20

60000/60000 [=====] - 3s 56us/step - loss: 0.3559
- acc: 0.8929 - val_loss: 0.2211 - val_acc: 0.9332

Epoch 5/20

60000/60000 [=====] - 3s 55us/step - loss: 0.3343
- acc: 0.8995 - val_loss: 0.2093 - val_acc: 0.9383

Epoch 6/20

60000/60000 [=====] - 3s 54us/step - loss: 0.3218
- acc: 0.9037 - val_loss: 0.2034 - val_acc: 0.9403

Epoch 7/20

60000/60000 [=====] - 3s 55us/step - loss: 0.3060
- acc: 0.9067 - val_loss: 0.1909 - val_acc: 0.9440

Epoch 8/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2969
- acc: 0.9093 - val_loss: 0.1818 - val_acc: 0.9460

Epoch 9/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2830
- acc: 0.9147 - val_loss: 0.1766 - val_acc: 0.9468

Epoch 10/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2684
- acc: 0.9189 - val_loss: 0.1699 - val_acc: 0.9491

Epoch 11/20

60000/60000 [=====] - 3s 57us/step - loss: 0.2599
- acc: 0.9216 - val_loss: 0.1631 - val_acc: 0.9502

Epoch 12/20

60000/60000 [=====] - 3s 57us/step - loss: 0.2501
- acc: 0.9237 - val_loss: 0.1567 - val_acc: 0.9530

Epoch 13/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2405
- acc: 0.9279 - val_loss: 0.1487 - val_acc: 0.9557

Epoch 14/20

60000/60000 [=====] - 3s 56us/step - loss: 0.2280
- acc: 0.9304 - val_loss: 0.1421 - val_acc: 0.9572

Epoch 15/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2165
- acc: 0.9338 - val_loss: 0.1322 - val_acc: 0.9600

Epoch 16/20

60000/60000 [=====] - 3s 56us/step - loss: 0.2107
- acc: 0.9369 - val_loss: 0.1310 - val_acc: 0.9592

Epoch 17/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2013
- acc: 0.9403 - val_loss: 0.1204 - val_acc: 0.9642

Epoch 18/20

60000/60000 [=====] - 3s 55us/step - loss: 0.1907
- acc: 0.9423 - val_loss: 0.1207 - val_acc: 0.9638

Epoch 19/20

60000/60000 [=====] - 3s 56us/step - loss: 0.1845
- acc: 0.9443 - val_loss: 0.1119 - val_acc: 0.9671

Epoch 20/20

60000/60000 [=====] - 3s 54us/step - loss: 0.1766
- acc: 0.9464 - val_loss: 0.1105 - val_acc: 0.9670

In [66]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.11045308456420898

Test accuracy: 0.967



In [67]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP +Batch-norm + 7 Dropout's + AdamOptimizer

In [68]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.7))
```

```
model_drop.add(Dense(346, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.63))
```

```
model_drop.add(Dense(496, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.57))
```

```
model_drop.add(Dense(639, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.19))
```

```
model_drop.add(Dense(99, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.87))
```

```
model_drop.add(Dense(750, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.1))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

WARNING:tensorflow:Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Model: "sequential_16"

Layer (type)	Output Shape	Param #
=====		
dense_61 (Dense)	(None, 512)	401920
batch_normalization_20 (Batch Normalization)	(None, 512)	2048
dropout_21 (Dropout)	(None, 512)	0
dense_62 (Dense)	(None, 128)	65664
batch_normalization_21 (Batch Normalization)	(None, 128)	512
dropout_22 (Dropout)	(None, 128)	0
dense_63 (Dense)	(None, 346)	44634
batch_normalization_22 (Batch Normalization)	(None, 346)	1384
dropout_23 (Dropout)	(None, 346)	0
dense_64 (Dense)	(None, 496)	172112
batch_normalization_23 (Batch Normalization)	(None, 496)	1984
dropout_24 (Dropout)	(None, 496)	0
dense_65 (Dense)	(None, 639)	317583
batch_normalization_24 (Batch Normalization)	(None, 639)	2556
dropout_25 (Dropout)	(None, 639)	0
dense_66 (Dense)	(None, 99)	63360
batch_normalization_25 (Batch Normalization)	(None, 99)	396
dropout_26 (Dropout)	(None, 99)	0
dense_67 (Dense)	(None, 750)	75000
batch_normalization_26 (Batch Normalization)	(None, 750)	3000
dropout_27 (Dropout)	(None, 750)	0
dense_68 (Dense)	(None, 10)	7510
=====		
Total params: 1,159,663		
Trainable params: 1,153,723		
Non-trainable params: 5,940		

In [69]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 181us/step - loss: 2.00

71 - acc: 0.2829 - val_loss: 0.9487 - val_acc: 0.6370

Epoch 2/20

60000/60000 [=====] - 7s 123us/step - loss: 1.067

1 - acc: 0.6132 - val_loss: 0.5183 - val_acc: 0.7925

Epoch 3/20

60000/60000 [=====] - 7s 122us/step - loss: 0.780

5 - acc: 0.7487 - val_loss: 0.3428 - val_acc: 0.9032

Epoch 4/20

60000/60000 [=====] - 7s 122us/step - loss: 0.617

2 - acc: 0.8172 - val_loss: 0.2718 - val_acc: 0.9270

Epoch 5/20

60000/60000 [=====] - 7s 122us/step - loss: 0.513

1 - acc: 0.8576 - val_loss: 0.2597 - val_acc: 0.9339

Epoch 6/20

60000/60000 [=====] - 7s 122us/step - loss: 0.469

3 - acc: 0.8768 - val_loss: 0.1974 - val_acc: 0.9508

Epoch 7/20

60000/60000 [=====] - 7s 122us/step - loss: 0.416

0 - acc: 0.8938 - val_loss: 0.1824 - val_acc: 0.9552

Epoch 8/20

60000/60000 [=====] - 7s 120us/step - loss: 0.384

8 - acc: 0.9035 - val_loss: 0.1783 - val_acc: 0.9594

Epoch 9/20

60000/60000 [=====] - 7s 121us/step - loss: 0.361

5 - acc: 0.9106 - val_loss: 0.1863 - val_acc: 0.9590

Epoch 10/20

60000/60000 [=====] - 7s 121us/step - loss: 0.342

1 - acc: 0.9140 - val_loss: 0.1516 - val_acc: 0.9649

Epoch 11/20

60000/60000 [=====] - 7s 122us/step - loss: 0.323

8 - acc: 0.9216 - val_loss: 0.1432 - val_acc: 0.9669

Epoch 12/20

60000/60000 [=====] - 7s 123us/step - loss: 0.309

3 - acc: 0.9256 - val_loss: 0.1470 - val_acc: 0.9661

Epoch 13/20

60000/60000 [=====] - 7s 122us/step - loss: 0.300

2 - acc: 0.9263 - val_loss: 0.1489 - val_acc: 0.9671

Epoch 14/20

60000/60000 [=====] - 7s 121us/step - loss: 0.294

3 - acc: 0.9298 - val_loss: 0.1334 - val_acc: 0.9700

Epoch 15/20

60000/60000 [=====] - 7s 122us/step - loss: 0.277

0 - acc: 0.9349 - val_loss: 0.1334 - val_acc: 0.9714

Epoch 16/20

60000/60000 [=====] - 7s 118us/step - loss: 0.261

8 - acc: 0.9384 - val_loss: 0.1340 - val_acc: 0.9713

Epoch 17/20

60000/60000 [=====] - 7s 118us/step - loss: 0.252

1 - acc: 0.9404 - val_loss: 0.1304 - val_acc: 0.9715

Epoch 18/20

60000/60000 [=====] - 7s 122us/step - loss: 0.245

6 - acc: 0.9422 - val_loss: 0.1152 - val_acc: 0.9742

Epoch 19/20

60000/60000 [=====] - 7s 123us/step - loss: 0.236

8 - acc: 0.9438 - val_loss: 0.1303 - val_acc: 0.9726

Epoch 20/20

60000/60000 [=====] - 7s 122us/step - loss: 0.232

2 - acc: 0.9459 - val_loss: 0.1185 - val_acc: 0.9742

In [70]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1184512525127735

Test accuracy: 0.9742



In [71]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[4].flatten().reshape(-1,1)
h5_w = w_after[4].flatten().reshape(-1,1)
h6_w = w_after[4].flatten().reshape(-1,1)
h7_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 8, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 8, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 8, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='y')
plt.xlabel('Hidden Layer 3 ')


plt.subplot(1, 8, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='c')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 8, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='m')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 8, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h6_w, color='g')
plt.xlabel('Hidden Layer 6 ')

plt.subplot(1, 8, 7)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h7_w, color='w')
plt.xlabel('Hidden Layer 7 ')

plt.subplot(1, 8, 8)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-norm+ 3 Dropout's + AdamOptimizer

In [72]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(624, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(370, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.4))
```

```
model_drop.add(Dense(82, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed = None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.3))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Model: "sequential_17"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_69 (Dense)	(None, 624)	489840
batch_normalization_27 (Batch Normalization)	(None, 624)	2496
dropout_28 (Dropout)	(None, 624)	0
dense_70 (Dense)	(None, 370)	231250
batch_normalization_28 (Batch Normalization)	(None, 370)	1480
dropout_29 (Dropout)	(None, 370)	0
dense_71 (Dense)	(None, 82)	30422
batch_normalization_29 (Batch Normalization)	(None, 82)	328
dropout_30 (Dropout)	(None, 82)	0
dense_72 (Dense)	(None, 10)	830
=====	=====	=====
Total params: 756,646		
Trainable params: 754,494		
Non-trainable params: 2,152		

In [73]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 116us/step - loss: 0.568

7 - acc: 0.8232 - val_loss: 0.2728 - val_acc: 0.9202

Epoch 2/20

60000/60000 [=====] - 4s 69us/step - loss: 0.3768

- acc: 0.8873 - val_loss: 0.2421 - val_acc: 0.9277

Epoch 3/20

60000/60000 [=====] - 4s 69us/step - loss: 0.3294

- acc: 0.9004 - val_loss: 0.2071 - val_acc: 0.9396

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.3011

- acc: 0.9099 - val_loss: 0.1889 - val_acc: 0.9442

Epoch 5/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2769

- acc: 0.9174 - val_loss: 0.1728 - val_acc: 0.9474

Epoch 6/20

60000/60000 [=====] - 4s 67us/step - loss: 0.2623

- acc: 0.9220 - val_loss: 0.1624 - val_acc: 0.9520

Epoch 7/20

60000/60000 [=====] - 4s 69us/step - loss: 0.2459

- acc: 0.9254 - val_loss: 0.1512 - val_acc: 0.9540

Epoch 8/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2297

- acc: 0.9312 - val_loss: 0.1445 - val_acc: 0.9558

Epoch 9/20

60000/60000 [=====] - 4s 69us/step - loss: 0.2220

- acc: 0.9339 - val_loss: 0.1330 - val_acc: 0.9600

Epoch 10/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2088

- acc: 0.9373 - val_loss: 0.1269 - val_acc: 0.9622

Epoch 11/20

60000/60000 [=====] - 4s 70us/step - loss: 0.1980

- acc: 0.9405 - val_loss: 0.1233 - val_acc: 0.9626

Epoch 12/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1934

- acc: 0.9421 - val_loss: 0.1189 - val_acc: 0.9639

Epoch 13/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1834

- acc: 0.9449 - val_loss: 0.1136 - val_acc: 0.9656

Epoch 14/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1790

- acc: 0.9465 - val_loss: 0.1120 - val_acc: 0.9663

Epoch 15/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1714

- acc: 0.9490 - val_loss: 0.1038 - val_acc: 0.9682

Epoch 16/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1597

- acc: 0.9521 - val_loss: 0.0981 - val_acc: 0.9708

Epoch 17/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1604

- acc: 0.9519 - val_loss: 0.0987 - val_acc: 0.9715

Epoch 18/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1491

- acc: 0.9545 - val_loss: 0.0934 - val_acc: 0.9739

Epoch 19/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1509

- acc: 0.9539 - val_loss: 0.0916 - val_acc: 0.9730

Epoch 20/20

60000/60000 [=====] - 4s 67us/step - loss: 0.1439

- acc: 0.9557 - val_loss: 0.0917 - val_acc: 0.9726

In [74]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09171288008186966

Test accuracy: 0.9726



In [75]:

```
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='w')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP +Batch-norm + 5 Dropout's + AdamOptimizer

In [76]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(710, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(623, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.7))
```

```
model_drop.add(Dense(548, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed = None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.8))
```

```
model_drop.add(Dense(475, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed = None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.2))
```

```
model_drop.add(Dense(316, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed = None)))
```

```
model_drop.add(BatchNormalization())
```

```
model_drop.add(Dropout(0.37))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Model: "sequential_18"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_73 (Dense)	(None, 710)	557350
batch_normalization_30 (Batch Normalization)	(None, 710)	2840
dropout_31 (Dropout)	(None, 710)	0
dense_74 (Dense)	(None, 623)	442953
batch_normalization_31 (Batch Normalization)	(None, 623)	2492
dropout_32 (Dropout)	(None, 623)	0
dense_75 (Dense)	(None, 548)	341952
batch_normalization_32 (Batch Normalization)	(None, 548)	2192
dropout_33 (Dropout)	(None, 548)	0
dense_76 (Dense)	(None, 475)	260775
batch_normalization_33 (Batch Normalization)	(None, 475)	1900
dropout_34 (Dropout)	(None, 475)	0
dense_77 (Dense)	(None, 316)	150416
batch_normalization_34 (Batch Normalization)	(None, 316)	1264
dropout_35 (Dropout)	(None, 316)	0
dense_78 (Dense)	(None, 10)	3170
=====	=====	=====
Total params: 1,767,304		
Trainable params: 1,761,960		
Non-trainable params: 5,344		

In [77]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 157us/step - loss: 0.983

8 - acc: 0.6901 - val_loss: 0.3823 - val_acc: 0.9008

Epoch 2/20

60000/60000 [=====] - 6s 94us/step - loss: 0.5394

- acc: 0.8332 - val_loss: 0.2944 - val_acc: 0.9163

Epoch 3/20

60000/60000 [=====] - 6s 92us/step - loss: 0.4524

- acc: 0.8621 - val_loss: 0.2613 - val_acc: 0.9254

Epoch 4/20

60000/60000 [=====] - 6s 95us/step - loss: 0.4109

- acc: 0.8760 - val_loss: 0.2539 - val_acc: 0.9267

Epoch 5/20

60000/60000 [=====] - 6s 95us/step - loss: 0.3879

- acc: 0.8828 - val_loss: 0.2362 - val_acc: 0.9329

Epoch 6/20

60000/60000 [=====] - 6s 94us/step - loss: 0.3642

- acc: 0.8904 - val_loss: 0.2155 - val_acc: 0.9381

Epoch 7/20

60000/60000 [=====] - 6s 94us/step - loss: 0.3487

- acc: 0.8970 - val_loss: 0.2056 - val_acc: 0.9408

Epoch 8/20

60000/60000 [=====] - 6s 93us/step - loss: 0.3305

- acc: 0.8998 - val_loss: 0.1960 - val_acc: 0.9451

Epoch 9/20

60000/60000 [=====] - 6s 96us/step - loss: 0.3190

- acc: 0.9041 - val_loss: 0.1848 - val_acc: 0.9496

Epoch 10/20

60000/60000 [=====] - 6s 94us/step - loss: 0.3058

- acc: 0.9087 - val_loss: 0.1767 - val_acc: 0.9499

Epoch 11/20

60000/60000 [=====] - 6s 93us/step - loss: 0.2935

- acc: 0.9121 - val_loss: 0.1722 - val_acc: 0.9513

Epoch 12/20

60000/60000 [=====] - 6s 92us/step - loss: 0.2866

- acc: 0.9145 - val_loss: 0.1586 - val_acc: 0.9545

Epoch 13/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2754

- acc: 0.9180 - val_loss: 0.1582 - val_acc: 0.9575

Epoch 14/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2673

- acc: 0.9205 - val_loss: 0.1513 - val_acc: 0.9579

Epoch 15/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2564

- acc: 0.9248 - val_loss: 0.1527 - val_acc: 0.9570

Epoch 16/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2507

- acc: 0.9265 - val_loss: 0.1433 - val_acc: 0.9616

Epoch 17/20

60000/60000 [=====] - 6s 93us/step - loss: 0.2377

- acc: 0.9298 - val_loss: 0.1433 - val_acc: 0.9630

Epoch 18/20

60000/60000 [=====] - 6s 92us/step - loss: 0.2283

- acc: 0.9319 - val_loss: 0.1380 - val_acc: 0.9638

Epoch 19/20

60000/60000 [=====] - 6s 93us/step - loss: 0.2253

- acc: 0.9349 - val_loss: 0.1301 - val_acc: 0.9639

Epoch 20/20

60000/60000 [=====] - 6s 92us/step - loss: 0.2188

- acc: 0.9347 - val_loss: 0.1274 - val_acc: 0.9651

In [78]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12741360827535392

Test accuracy: 0.9651



In [79]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(20,5))
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1,6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='c')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='m')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [0]:

```

from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer
=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, st
ddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='ada
m')

    return model

```

In [0]:

```

# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-p
ython-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batc
h_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs =-1)
grid_result = grid.fit(X_train, Y_train)

```

In [82]:

```

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

```

```

Best: 0.977383 using {'activ': 'sigmoid'}
0.977383 (0.001965) with: {'activ': 'sigmoid'}
0.976550 (0.003953) with: {'activ': 'relu'}

```

In [98]:

```
# http://zetcode.com/python/prettytable/
from prettytable import PrettyTable

ptable = PrettyTable()

ptable.field_names = [ "Model", "Accuracy"]
ptable.add_row(["softmax",90.87])
ptable.add_row(["MLP+Sigmoid+SGDoptimizer",87.71])
ptable.add_row(["MLP+Sigmoid+ADAM",98.06])
ptable.add_row(["MLP+SGD+ReLU",96.27])
ptable.add_row(["MLP+ReLU+ADAM",97.6])
print(ptable)

ptable = PrettyTable()
ptable.field_names =["Model","Layers","Accuracy"]
ptable.add_row(["MLP+Batch-norm+ADAM",2,97.41])
ptable.add_row(["MLP+Batch-norm+ADAM",3,97.36])
ptable.add_row(["MLP+Batch-norm+ADAM",5,97.66])
ptable.add_row(["MLP+Batch-norm+ADAM",7,97.6])
print(ptable)

ptable = PrettyTable()
ptable.field_names= ["Model","dropouts","Accuracy"]
ptable.add_row(["MLP+dropouts+ADAM",2,97.89])
ptable.add_row(["MLP+Dropouts+ADAM",3,97.47])
ptable.add_row(["MLP+Dropouts+adam",5,97.4])
print(ptable)

ptable = PrettyTable()
ptable.field_names= ["Model","dropouts","layers","Accuracy"]
ptable.add_row(["MLP+Batch-norm+ADAM+DROPOUTS",2,2,96.7])
ptable.add_row(["MLP+Batch-norm+ADAM+DROPOUTS",7,7,97.42])
ptable.add_row(["MLP+Batch-norm+ADAM", 3,3,97.26 ])
ptable.add_row(["MLP+Batch-norm+ADAM+DROPOUTS", 5,5,96.51])
print(ptable)
```

Model	Accuracy
softmax	90.87
MLP+Sigmoid+SGDoptimizer	87.71
MLP+Sigmoid+ADAM	98.06
MLP+SGD+ReLU	96.27
MLP+ReLU+ADAM	97.6

Model	Layers	Accuracy
MLP+Batch-norm+ADAM	2	97.41
MLP+Batch-norm+ADAM	3	97.36
MLP+Batch-norm+ADAM	5	97.66
MLP+Batch-norm+ADAM	7	97.6

Model	dropouts	Accuracy
MLP+dropouts+ADAM	2	97.89
MLP+Dropouts+ADAM	3	97.47
MLP+Dropouts+adam	5	97.4

Model	dropouts	layers	Accuracy
MLP+Batch-norm+ADAM+DROPOUTS	2	2	96.7
MLP+Batch-norm+ADAM+DROPOUTS	7	7	97.42
MLP+Batch-norm+ADAM	3	3	97.26
MLP+Batch-norm+ADAM+DROPOUTS	5	5	96.51